SSGAC/CESR Social-Science Genetics Methods Workshop

Statistical genetics on All of Us

Session 1: a phenotype, a polygenic index, and a regression on v9 data

Jeffrey M. Swigert, PhD

Center for Economic and Social Research, University of Southern California

Fall 2026

Generously supported by NIH/NIA Grant R24-AG065184 and the All of Us Research Program

Do this now: start your cloud environment

In the workshop workspace: **Apps** tab → **+ New app instance** → **JupyterLab**. Creation takes a few minutes; we continue while it boots.

Setting	Choose	Why
Environment	Standard VM	The Spark cluster option is for Hail.
Machine	`n1-standard-8`: 8 CPUs, 30 GB memory	The first machine in the menu. Enough for one chromosome.
Disk	100 GB	Today uses under 5 GB. Disks grow, never shrink.
GPU	None	Nothing today uses a GPU.
Autostop	On, 1 hour	Stops the VM after 1 idle hour: a clean stop when the session is over.

About $0.41 per hour while running; a stopped VM bills only for its disk.

What we will do today

In this session: seven steps, each one script.

Step	What we do	Script (in `scripts/`)
0	Get the code, fill in settings	`00_setup.sh`
1	Put chromosome-22 genotypes and ancestry data on the VM	`01_fetch_genotypes.sh`
2	Build a height phenotype from the CDR	`02_build_phenotype.sh`
3	Summarize it: a summary table and two figures	`03_explore_phenotype.sh`
4	Build a polygenic index (PGI); check its histogram	`04_build_pgi.sh`
5	Regress height on the PGI and covariates	`05_pgi_regression.sh`
6	Save the run; download aggregate results	`06_save_run.sh`

For each step: a slide on what it does, then a **Run it** slide with the commands. We run each step together before moving on.

The repository has the details

`github.com/jeffreyswigert/aou-statgen-lab`

Open it in a browser tab now and keep it open. The slides are brief on purpose; the repository explains every step in more detail:

- **README**: what each step does, why, and where the data and weights come from.
- **The scripts**: every line has a comment saying what it does, why, and what it looks like in your own project.
- `COMMANDS.txt`: every command from today, in order, with a note on each. Copy and paste from it.
- `cheatsheet.pdf` and `slides.pdf`: for after class.

If a step goes by too fast today, its script and the README are where the explanation lives.

All of Us data use rules

Rule	In this session
Participant-level data stays in the Workbench: no copying, downloading, or screenshots.	Person-level files live in `lab_data/` and `work/` only.
No count of 1–20 participants may be shared, directly or by calculation.	Step 6 checks `results/` before anything is saved or downloaded.
Downloads are monitored; downloading participant-level data triggers an alert.	Downloads come from `results/` only.
No participant data or counts under 20 in code on GitHub.	`.gitignore` excludes data and output folders.
Notify the program 2 weeks before a publication or presentation.	For your own projects.

GitHub: what it is and how we use it

Git records versions of the files in a folder; each saved version is a *commit*. **GitHub** hosts a copy of the folder (a *repository*) online.

Today it does three things:

1. **Gets the code onto the VM**: `git clone` once, `git pull` for updates.
2. **Records which code produced a result**: step 6 writes the commit ID into `results/run_manifest.txt`.
3. **Runs identical code** in every participant's workspace.

GitHub is outside the Workbench. A repository holds code, settings templates, and documentation. It should never hold data, results, filled-in settings, or notebook outputs.

`.gitignore`: keeping data out of GitHub

A `.gitignore` file lists files and folders that git never tracks. This repository's:

```
config.sh              # your filled-in settings
lab_data/              # copied-in data (person-level)
work/                  # intermediate files (person-level)
results/               # aggregate outputs (reviewed and downloaded by hand)
runs/                  # saved archives
.ipynb_checkpoints/    # notebook autosaves
```

For any project: ignore data and output folders by name; commit a settings template and ignore the filled-in copy; clear notebook outputs before committing.

More practices in the README, "GitHub in this workflow."

Run it – Step 0: get the code and fill in settings

JupyterLab Launcher → **Terminal**. One command at a time (or paste from `COMMANDS.txt`):

```
git clone https://github.com/jeffreyswigert/aou-statgen-lab.git   # copy the repository here
cd aou-statgen-lab                  # move into it
cp config.example.sh config.sh      # your own copy of the settings
nano config.sh                      # opens the settings file: fill in the two values below
```

In nano, fill in the two values (shown on screen; also on the workspace page: Resources tab → bucket; Overview → Google project):

```
export WORKSHOP_BUCKET="gs://..."    # the workshop workspace's bucket
export BILLING_PROJECT="..."         # the Google project billed for queries and copies
```

Save and exit nano: `Ctrl-O`, `Enter`, `Ctrl-X`. If that does not work in your browser (Windows): `Ctrl-X`, `Y`, `Enter`.

Only then, back at the prompt:

```
bash scripts/00_setup.sh            # check tools, settings, disk
```

Expected: the last line reads `Setup check passed.`

Where today's data are

1. **The CDR**, in BigQuery: `wb-silky-artichoke-2408.C2025Q4R6`

   Tables `measurement`, `person`, `concept` → the height phenotype (step 2)

2. **The All of Us controlled bucket**: `gs://vwb-aou-datasets-controlled/v9/`

   `wgs/short_read/snpindel/aux/ancestry/ancestry_preds.tsv`: predicted ancestry and 16 PCs per person (step 1)

   `wgs/short_read/snpindel/acaf_threshold/`: whole-genome genotypes, one file set per chromosome (next slide)

3. **The workshop bucket**: `$WORKSHOP_BUCKET`

   `genotypes/chr22_hm3.{bed,bim,fam}`: the chromosome-22 subset, ready made (step 1)

   `pgi_workshop/height_weights.txt`: height PGI weights, Yengo et al. (2022) (step 4)

4. **The repository**: `data/hm3_chr22_hg38.tsv`, the public HapMap3 variant list

The controlled bucket is **requester-pays**: each copy names who pays (`--billing-project`).

Where the genotypes come from

- All of Us publishes whole-genome genotypes (the ACAF callset) per chromosome: for chromosome 22, 25 GB (PLINK 2) or 276 GB (PLINK 1). No HapMap3 subset.
- **HapMap3 variants**: about 1.2 million common variants that most GWAS report; the standard set for PGIs.
- So a PGI project starts by keeping those variants and renaming them to rsIDs (`scripts/subset_aou_genotypes.sh`):

```
# keep the HapMap3 variants and the chosen people; name variants chr22:pos:ref:alt
plink2 --pfile acaf --set-all-var-ids 'chr@:#:$r:$a' \
       --extract extract.txt --keep keep.txt --make-bed --out hm3_aouids
# rename the variants to rsIDs, to match the PGI weights
plink2 --bfile hm3_aouids --update-name rename.txt --make-bed --out chr22_hm3
```

I ran this before class; step 1 copies the result. (Building it yourself: "On your own" slides, and the README.)

Run it – Step 1: the genotypes and the ancestry file

```
bash scripts/01_fetch_genotypes.sh      # under a minute; start it, then read on
```

What the script does, line by line (`scripts/01_fetch_genotypes.sh`):

Line	What it does	In your own project
`source scripts/common.sh`	loads config.sh, the safety switches, the billing flag	one shared setup file
`mkdir -p lab_data`	a folder for copied-in data	data outside the code folder
`for ext in bed bim fam; do ... done`	runs the copy three times, once per extension	files that travel together
`gcloud ... storage cp SRC DEST`	copies one file from a bucket to disk	the same command for any `gs://` file
`gcloud ... cp "$ANC_SRC" lab_data/...`	the ancestry file, from the requester-pays bucket	`-billing-project` names who pays
`ls -lh ...; wc -l < FILE`	file sizes; lines per file	check every copy

Run it – Step 1: look at what arrived

When the script finishes:

```
head -n 2 lab_data/chr22_hm3.fam                        # first 2 lines: one line per person
head -n 2 lab_data/chr22_hm3.bim                        # one line per variant
head -n 1 lab_data/ancestry_preds.tsv | tr '\t' '\n'  # the header, one column name per line
```

Piece	Meaning
`head -n N FILE`	print the first N lines of FILE (the way to look at any text file)
`A | B`	a *pipe*: what A prints goes into B instead of onto the screen
`tr '\t' '\n'`	replace every tab with a line break, so a wide header reads as a list

Expected: the script's last lines read `chr22: ... variants staged`. Everyone's count is the same.

The files we just copied

File (one line =)	Columns	Note
`chr22_hm3.fam` (person)	FID, IID, father, mother, sex, phenotype	FID is 0 for everyone. IID is the person ID.
`chr22_hm3.bim` (variant)	chromosome, ID, cM, position, allele 1, allele 2	ID is an rsID; positions are GRCh38.
`chr22_hm3.bed`	genotypes, binary	Never open as text.
`ancestry_preds.tsv` (person)	`research_id`, `ancestry_pred`, `probabilities`, `pca_features`, `ancestry_pred_other`	`research_id` = CDR `person_id`. `pca_features`: 16 PCs in one text field.

Made-up lines. One person ID, three column names: `IID`, `person_id`, `research_id`.

`chr22_hm3.fam`

```
0 1000010 0 0 2 -9
```

`chr22_hm3.bim`

```
22 rs1234567 0 17054321 A
    G
```

`ancestry_preds.tsv`

```
1000010 eur [0.01,..] [-0.009,..] eur
```

Where phenotype data live in All of Us

Source	What it holds	CDR table(s), and how rows are labeled
Program measurements	Height, weight, BMI, waist and hip, blood pressure, heart rate; taken at enrollment, some self-reported remotely	`measurement`; `measurement_source_concept_id` in the All of Us (PPI) vocabulary, e.g. 903133
Surveys	The Basics, Overall Health, Lifestyle, health history, and more	`observation` (also `ds_survey`); question = `observation_source_concept_id`, answer = `value_source_concept_id`
Electronic health records	Diagnoses, medications, procedures, lab results and vitals, visits; from participating sites	`condition_occurrence`, `drug_exposure`, `procedure_occurrence`, `measurement`, `visit_occurrence`; standard OMOP concepts, e.g. height 3036277
Wearables (Fitbit)	Daily activity, heart rate, sleep	`activity_summary`, `heart_rate_summary`, `sleep_daily_summary`, ...; one row per person-day or minute
Demographics	Year of birth, sex at birth, race, ethnicity	`person`

Height is in `measurement` twice: program-measured (903133) and from health records (3036277).

The Data Browser: finding a variable before you query

`databrowser.researchallofus.org`: public, no login; participant counts rounded to the nearest 20.

Question	Where	What you get
Which survey asks about X, and how was it worded?	Surveys	the question text, answer options, participants per answer, concept IDs
How many participants have a diagnosis, drug, lab, or procedure?	EHR domains	concept name and ID, participant count; for labs, the value distribution
What does a program measurement look like?	Physical Measurements	distributions of height, weight, blood pressure, ... by sex and age
Is a variant in the data?	Genomic Variants	search by gene, rsID, or position; allele counts and frequencies

The concept ID on each card goes in a query, or in a row of `data/phenotypes.tsv`. The same lookup in SQL:

```
SELECT concept_id, concept_name FROM `$CDR_DATASET.concept`
WHERE LOWER(concept_name) LIKE '%height%'
```

Height: the concept choice

- Height is in the CDR `measurement` table: one row per measurement, often several rows per person.
- We select `measurement_source_concept_id = 903133`: height measured by the program at enrollment, plus remote self-report.
- Not the standard concept 3036277: it also returns heights from electronic health records.
- Year of birth and sex at birth come from the `person` table.

The query and the files it writes

The first of two queries (`scripts/02_build_phenotype.sh`):

```
SELECT person_id, value_as_number          -- which columns
FROM `$CDR_DATASET.measurement`            -- which table
WHERE measurement_source_concept_id IN (903133)   -- which rows: the concept,
  AND value_as_number IS NOT NULL                 --   with a value
```

The second gets `year_of_birth` and `sex_at_birth` from `person`. Results go to `work/` (made-up values):

`work/raw_pheno.csv`

```
person_id,value_as_number
1000010,172.4
1000010,172.0
1000020,158.9
```

`work/raw_person.csv`

```
person_id,year_of_birth,sex_at_birth
1000010,1961,Female
1000020,1983,Female
```

Person 1000010 has two rows in `raw_pheno.csv`: the next step reduces them to one.

Run it – Step 2: build the phenotype

```
DRY_RUN=1 bash scripts/02_build_phenotype.sh    # prints the two queries; runs nothing
bash scripts/02_build_phenotype.sh              # runs the queries, then cleans (about 1 min)
```

The cleaning (`scripts/build_pheno.py`, lightly simplified):

```
value = statistics.median(values)                    # 1: each person's median
if value < LO or value > HI: continue                # 2: drop outside 100-250 cm
if sex not in ('Male', 'Female'): sex = 'other'      # 3: sex at birth
age = 2026 - int(yob)                                # 4: age
z = (value - mu) / sd                                # 5: standardize within sex
```

Output `work/aou_pheno.tsv`, one row per person: `person_id value value_z age sex`.

Expected (the script prints the sample at each step):

```
measurement records pulled          ...
removed: outside 100-250 cm         <=20
people in analysis file             ...
```

Exploratory data analysis: the summary table

Input: `work/aou_pheno.tsv` (columns `person_id value value_z age sex`). Check it from the terminal; output stays on your screen:

```
head -n 3 work/aou_pheno.tsv | column -t       # first 3 lines, columns lined up
wc -l work/aou_pheno.tsv                       # count lines: rows + 1 header line
cut -f 5 work/aou_pheno.tsv | sort | uniq -c   # keep column 5, sort it, count each value
```

Outputs in `results/`: `aou_pheno_summary.txt` (the table), `aou_pheno_hist.png`, `aou_pheno_kde_sex.png`.

The table has one row per group: Obs, Mean, Std. dev., p25 / Median / p75. Percentiles stand in for Min and Max, because each of those is one person's value.

Code: `scripts/pheno_eda.py - summary_row()` makes each table row

Run it – Step 3: the summary table and figures

```
bash scripts/03_explore_phenotype.sh
cat results/aou_pheno_summary.txt
```

Expected layout (made-up numbers):

```
     Variable |       Obs        Mean    Std. dev.       p25      Median       p75
--------------+------------------------------------------------------------------
       height |       ...      169.10         9.80     162.0       168.9     176.2
         Male |       ...      175.90         7.40     171.0       175.9     180.8
       Female |       ...      162.40         6.90     157.8       162.3     167.0
```

Check: US adult means are about 176 cm (men) and 162 cm (women); SDs 7–8 cm.

Open the figures by double-clicking them in the JupyterLab file browser, folder `results/`.

Building the PGI

For person i, the PGI adds up, over variants j, the variant's weight times the person's number of copies of the effect allele:

$$\mathrm{PGI}_i = \sum_j w_j \, g_{ij}, \qquad g_{ij} \in \{0, 1, 2\}$$

Weights w_j: height, Yengo et al. (2022), the GIANT GWAS of 5.3 million people; PGS Catalog score PGS002802.

File	Contents
input `lab_data/chr22_hm3.{bed,bim,fam}`	genotypes
input `lab_data/height_weights.txt`	weights: `rsid effect_allele weight`
output `work/aou_pgi.sscore`	one row per person; `SCORE1_SUM` is the PGI
output `work/aou_pgi.log`	PLINK's log, including the match count

Where a genotype is missing, PLINK uses the variant's average allele count.

The PLINK command

From `scripts/04_build_pgi.sh` (with `CHROM=22`):

```
p2 --bfile lab_data/chr22_hm3 \
   --score lab_data/height_weights.txt 1 2 3 header cols=+scoresums list-variants \
   --out work/aou_pgi
```

Part	Meaning
`p2`	`plink2` with our CPU and memory settings (`scripts/common.sh`)
`-bfile lab_data/chr22_hm3`	read `chr22_hm3.bed`, `.bim`, `.fam`
`-score FILE 1 2 3`	weights file; variant ID in column 1, effect allele in 2, weight in 3
`header`	the first line of the weights file is column names
`cols=+scoresums`	also write each person's sum (`SCORE1_SUM`)
`-out work/aou_pgi`	prefix for the output files

Run it – Step 4: build the PGI

```
bash scripts/04_build_pgi.sh                              # weights, PLINK, QC histogram
grep -E 'variants processed|skipped' work/aou_pgi.log     # the match count, from the log
wc -l < lab_data/height_weights.txt                       # weight rows + 1 header line
```

The match count. `variants processed` = weight rows that found their variant in the genotype file. PLINK skips unmatched rows without an error.

- Close to the number of weight rows: expected.
- Near zero: the variant IDs do not match (rsIDs vs. `chr22:pos:ref:alt`), or the column numbers are wrong.

QC on the PGI: a histogram

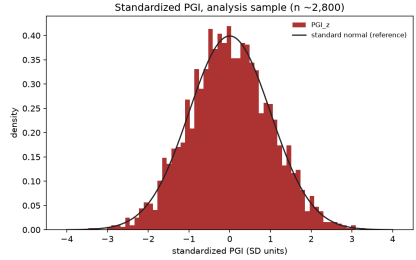
Step 4 standardizes the PGI and plots its histogram with a normal curve for reference:

$$\text{PGI_z}_i = \frac{\text{PGI}_i - \text{mean(PGI)}}{\text{SD(PGI)}}$$

```
pgi_z = (raw - raw.mean()) / raw.std()
ax.hist(pgi_z, bins=60, density=True)
```

Expected: a peak just below 0 and a longer right tail than the curve; skewness about $+0.6$ (`results/aou_pgi_summary.txt`).

Code: `scripts/pgi_eda.py`



Example output of this code on simulated data (not All of Us data).

Allele frequencies differ between genetic ancestries, so the real PGI is not one bell curve: one reason the regression adjusts for principal components.

The regression

$$\text{height_z}_i = \alpha + \beta\,\text{PGI_z}_i + \gamma\,\text{age}_i + \delta\,\text{female}_i + \sum_{k=1}^{5}\theta_k\,\text{PC}_{ik} + \varepsilon_i$$

Variable	Definition	File (step)
height_z	(height − mean for that sex) / SD for that sex	`work/aou_pheno.tsv` (2)
PGI_z	(PGI − mean) / SD, over the people in the model	`work/aou_pgi.sscore` (4)
age	2026 − year of birth	`work/aou_pheno.tsv` (2)
female	1 if Female, 0 if Male	`work/aou_pheno.tsv` (2)
PC_1–PC_5	first 5 genetic principal components	`lab_data/ancestry_preds.tsv` (1)

Estimated by ordinary least squares (OLS).

What the regression reports

- β: both variables are in SD units, so β is the difference in height, in SDs, per 1 SD of PGI, holding age, sex, and the PCs fixed.
- **The PCs** summarize broad genetic ancestry, so ancestry-related differences in height are not attributed to the PGI.
- **Incremental** $R^2 = R^2$ with $\mathrm{PGI_z} - R^2$ without it: the share of height variance the PGI explains beyond the covariates.
- **Scope**: one chromosome, and OLS, which treats participants as unrelated. A full PGI uses all 22 chromosomes.

Run it – Step 5: the regression

```
bash scripts/05_pgi_regression.sh
cat results/aou_pgi_regression.txt
```

The fit (scripts/pgi_regression.py):

```
X[:, 0] = (X[:, 0] - X[:, 0].mean()) / X[:, 0].std()   # PGI -> PGI_z
beta, se, r2_full = fit(X)          # height_z on PGI_z, age, female, PC1-PC5
_, _, r2_cov = fit(X[:, 1:])        # the same model without PGI_z
```

Expected layout (made-up numbers):

```
term          estimate       SE       t
PGI_z           0.0XXX   0.00XX    X.XX
...
R2 with PGI    = 0.0XXX
R2 covariates  = 0.0XXX
incremental R2 = 0.00XX
```

Saving results and downloading them

Folder	Contents	Can it leave the Workbench?
`lab_data/`	copied-in genotypes, ancestry file, weights	No
`work/`	person-level files: phenotype, PGI scores	No
`results/`	summary tables (`.txt`), figures (`.png`)	Yes, after the check and your review
`runs/`	archive of the run (copied to the bucket)	No

To download a file: open it and read it; then JupyterLab file browser → right-click → **Download**.

Figures and free text are not checked by the script: review them yourself. Copying out with `gcloud storage cp` is blocked; a large summary file needs an approved exception first.

What step 6 does

`scripts/06_save_run.sh`:

1. **Checks `results/`** for counts of 1–20 (direct, or from the difference of two counts) and for person-level files; then checks the counts the steps printed. If it finds any, it stops and names the file.
2. **Writes `results/run_manifest.txt`**: the git commit ID, the settings, the PLINK version, and checksums of the input files.
3. **Packs** `config.sh`, `scripts/`, `work/`, and `results/` into one archive in `runs/`.
4. **Copies** the archive to `$WORKSHOP_BUCKET/statgen-lab/` and lists it to confirm.

The archive holds person-level files, so it stays in the workspace bucket. It is never downloaded.

Code: `scripts/check_disclosure.py; scripts/check_counts.py`

Run it – Step 6: save, download, stop

```
bash scripts/06_save_run.sh
cat results/disclosure_report.txt                    # what the check found
grep code_version results/run_manifest.txt           # the commit that produced this run
gcloud storage ls "$WORKSHOP_BUCKET/statgen-lab/"    # list the bucket folder: the archive is
    there
```

Expected: `none found` twice, the printed-counts check with no pairs, then the archive's bucket path. If the check stops the save, it prints the file and the count: fix that cell and re-run.

Download one file: file browser → `results/` → right-click `aou_pgi_hist.png` → Download.

Stop the VM: Apps tab → your JupyterLab app → **Stop**. Stop keeps the app and its disk.

All the commands, in order

```
git clone https://github.com/jeffreyswigert/aou-statgen-lab.git && cd aou-statgen-lab
cp config.example.sh config.sh && nano config.sh   # WORKSHOP_BUCKET, BILLING_PROJECT
bash scripts/00_setup.sh                # check tools and settings
bash scripts/01_fetch_genotypes.sh      # genotypes + ancestry file       -> lab_data/
bash scripts/02_build_phenotype.sh      # CDR queries, one row per person -> work/
bash scripts/03_explore_phenotype.sh    # summary table, 2 figures        -> results/
bash scripts/04_build_pgi.sh            # PGI, match count, histogram     -> work/, results/
bash scripts/05_pgi_regression.sh       # the regression                  -> results/
bash scripts/06_save_run.sh             # check, manifest, archive to the workspace bucket
```

Settings to change (in `config.sh`, or before one command): `PHENO=ldl`, another phenotype (`bash scripts/02_build_phenotype.sh list`); `SAMPLE_MOD=10`, 1 person in 10 (a quick trial run); `GENO_SOURCE=aou`, build the subset from All of Us's files; `CHROM=21`, another chromosome.

After this slide: reference slides on autostop, setting up your own workspace, building the genotype subset, and the acronyms used today.

Thank you

Questions?

Code, README, commands, cheat sheet, slides:
`github.com/jeffreyswigert/aou-statgen-lab`

SSGAC: `https://www.thessgac.org`

Jeff Swigert: `js_850@usc.edu`

SSGAC/CESR Social-Science Genetics Methods Workshop. Generously supported by NIH/NIA Grant R24-AG065184 and the All of Us Research Program.

Reference: autostop, when to change it

Autostop stops the VM when it has been idle. The VM counts as **idle** when, for 15 minutes,

- average CPU use is below 50%, and
- no browser tab or terminal connection is active.

Keep autostop on (1 hour today) for interactive work.

Set it longer than the job (up to 14 days) when a job runs for hours with low CPU use:

- a single-threaded program on an 8-CPU VM (1 busy CPU = 12.5% use);
- a large copy from a bucket, which waits on the network, not the CPU;
- a script that waits on a long BigQuery query.

Otherwise the VM can stop in the middle of the job. **Turn autostop off** only if you will stop the VM yourself.

On your own: set up a workspace and add resources

Today the workspace, its data, and its bucket were set up for you. For your own project, in the Researcher Workbench:

1. **Create a workspace.** Give it an ID, a region, and the billing **pod** that pays for it. Genotypes need Controlled Tier access.
2. **Add the All of Us data.** Add the Controlled Tier v9 data collection and select **all resources**. This adds references to the CDR and the controlled bucket, not copies.
3. **Create a bucket** for your own files (a new Cloud Storage bucket resource). Its `gs://` address is `WORKSHOP_BUCKET` in `config.sh`.
4. **Find the Google project**: workspace details → **Google project** (it looks like `wb-word-word-1234`). It is `BILLING_PROJECT` in `config.sh`.
5. **Create the app** (slide 2).

On the VM, the added data appear as folders under `~/workspace/`:

```
ls ~/workspace/                                  # one folder per added resource
ls ~/workspace/vwb-aou-datasets-controlled-v9/   # shows: pooled  v7  v8  v9
```

On your own: the genotype subset, step by step

Today step 1 copied a subset built before class from All of Us's own files. To build one yourself:

	Step	For chromosome 22
1	Find All of Us's genotype files for the chromosome	`.../acaf_threshold/pgen/`: `.pgen` (genotypes, 25 GB), `.pvar` (variants), `.psam` (people)
2	Get a list of the variants you want, with GRCh38 positions	HapMap3 list, `data/hm3_chr22_hg38.tsv`: 15,414 variants
3	Match the list to All of Us's variants by position and alleles	All of Us names variants `chr22:pos:ref:alt`, not by rsID; 9,944 matched
4	Keep the matched variants (and the people you want)	first `plink2` command: `--extract`, `--keep`, `--make-bed`
5	Rename the variants to rsIDs	second `plink2` command: `--update-name`
6	Check the result	count the lines of the `.bim`; look at its first rows

The script skips sites where All of Us lists more than one alternate allele: 9,944 of 15,414 matched. The next slide runs all six steps with one command.

On your own: build the subset from All of Us

The same step 1, told to build the subset instead of copying it:

```
GENO_SOURCE=aou bash scripts/01_fetch_genotypes.sh    # about 6 minutes for chromosome 22
```

What it does (`scripts/subset_aou_genotypes.sh`):

1. Copies All of Us's two small chromosome-22 files: variants (`.pvar`) and people (`.psam`).
2. Matches the HapMap3 list (`data/hm3_chr22_hg38.tsv`) to the variants by position and alleles.
3. Copies the 25 GB genotype file (`.pgen`) to the VM: about 5 minutes, 25 GB of free disk.
4. Runs the two PLINK commands from "Where the genotypes come from": about 1 minute.
5. Deletes the 25 GB copy. The result is the same `lab_data/chr22_hm3` files as today's.

Other chromosomes: add `CHROM=21` and a HapMap3 list for it (`HM3_LIST=...`; see the README).

Autostop: a long copy can look idle. Keep the browser tab open, or set autostop longer than the job.

On your own: check the subset and keep it

After the build, check it before using it:

```
wc -l < lab_data/chr22_hm3.bim          # number of variants (9,944 for chromosome 22)
head -n 2 lab_data/chr22_hm3.bim        # the names should be rsIDs, not chr22:pos:ref:alt
ls -lh lab_data/chr22_hm3.*             # .bed is about 1.3 GB with everyone in it
```

Then store it in your workspace bucket, so the build runs once and not on every new VM:

```
for ext in bed bim fam; do
  gcloud storage cp lab_data/chr22_hm3.$ext "$WORKSHOP_BUCKET/genotypes/"
done
```

That is what was done for today: `scripts/prep/make_hm3_subset.sh` builds the subset and copies it to the bucket, and step 1 copies it back in seconds.

The files stay in the Workbench. The subset is person-level genotype data: it can go to a workspace bucket, never to your own computer or to GitHub.

Acronyms used today: the data

Acronym	Stands for	In this session
AoU	All of Us	the research program and its data
CDR	Curated Data Repository	the All of Us database the queries run against (v9 today)
OMOP	Observational Medical Outcomes Partnership	the CDR's table layout (`measurement`, `person`, ...)
ETL	Extract, Transform, Load	the program's pipeline from raw records into those tables; it assigns the standard codes
EHR	Electronic Health Record	hospital data; LDL and HDL come only from here
PPI	Participant-Provided Information	All of Us's own codes; height 903133, in `measurement_source_concept_id`
LOINC	Logical Observation Identifiers Names and Codes	standard codes for labs and measurements; height 3036277, in `measurement_concept_id`
SQL	Structured Query Language	the language of the two queries in step 2
CSV/TSV	Comma-/Tab-Separated Values	the plain-text tables in `work/`, `results/`, `data/`
VM	Virtual Machine	the cloud computer the Workbench starts for you

Acronyms used today: the analysis

Acronym	Stands for	In this session
GWAS	Genome-Wide Association Study	the study that estimated the per-variant weights
PGI, PGS	Polygenic Index, Polygenic Score	the weighted sum of genotypes from step 4; height weights: PGS Catalog PGS002802
HM3	HapMap 3	about a million common variants; the chr22 subset keeps only these
ACAF	Allele Count, Allele Frequency	the All of Us genotype release filtered on those two; source of the subset
rsID, SNP	Reference SNP ID; Single-Nucleotide Polymorphism	a variant's name, e.g. rs12345, as used in `.bim` and the weights file
FID, IID	Family ID, Individual ID	the two PLINK ID columns; All of Us sets FID to 0, so the scripts join on IID
QC	Quality Control	plausibility bounds, the match count, the histogram
PC	Principal Component	genetic ancestry summaries; PC1–PC5 are regression covariates
OLS	Ordinary Least Squares	the regression in step 5
SD	Standard Deviation	the unit of standardized height and of the PGI